

---

EECS 182      Deep Neural Networks  
 Fall 2025      Anant Sahai and Gireeja Ranade      Discussion 9

---

## 1. SSM Convolution Kernel

Consider a discrete-time linear time-invariant State-Space Model (SSM) of the form

$$\mathbf{x}_{k+1} = A\mathbf{x}_k + Bu_k, \quad (1)$$

$$y_k = C\mathbf{x}_k + Du_k, \quad (2)$$

For simplicity (and realism vis-a-vis modern state-space models), consider both  $u$  and  $y$  to be scalars while the state  $\mathbf{x}$  is a vector.

- (a) **Convolution Kernel and the Output Equation.** Given that the sequence length is  $L$  (input:  $(u_1, \dots, u_L)$ , output:  $(y_1, \dots, y_L)$ ) and assuming the initial state  $\mathbf{x}_0 = \mathbf{0}$ , show that the output  $y_k$  can be expressed as a *convolution* of the input sequence  $\{u_\ell\}_1^L$  with a kernel  $K = \{K_\ell\}_1^L$ :

$$y_k = \sum_{\ell=0}^L K_\ell u_{k-\ell},$$

where any  $u_{\leq 0}$  with a negative index is set to 0 (zero-padding). Also, find  $K$ .

- (b) **Efficient Computation with Convolutions.** These facts will be useful for this question:

- The convolution theorem states that if we have two sequences  $u[k]$  and  $v[k]$  where their discrete-time Fourier transforms are  $U(f)$  and  $V(f)$  respectively, then the discrete-time Fourier transform of their convolution  $W[f] = \mathcal{F}\{u[k] * v[k]\} = U(f)V(f)$ .
- The Fast Fourier Transform (FFT) algorithm can compute the discrete Fourier transform of a length- $N$  sequence in  $O(N \log N)$  time.
- The inverse discrete Fourier transform can also be computed using FFT in  $O(N \log N)$  time.

If we already know the kernel  $K$ , how much can we parallelize the computation of a scalar output sequence  $\{y_k\}$  for a scalar input sequence  $\{u_k\}$  of length  $L$ ? What is the minimum critical path length of the parallelized computation?

What about a naive, direct computation of  $y_k$  from the recursion?

- (c) **Efficient Kernel Computation.** Given  $A, B, C$  are matrices, how can we compute the kernel  $K$  efficiently? What are some strategies to parallelize kernel computation? You may assume  $L = 2^N$  for some  $N$  for simplicity.
- (d) **Adding structure to  $A$ .** Suppose  $A$  is a diagonal matrix where we represent each diagonal entry as  $a = \exp(\lambda_R + j\lambda_I)$ . How can we leverage this structure to compute the kernel  $K$  more efficiently?
- (e) **Linear Time-varying SSM.** Now consider a linear time-varying SSM of the form

$$\mathbf{x}_{k+1} = A_k \mathbf{x}_k + B_k u_k, \quad (3)$$

$$y_k = C_k \mathbf{x}_k + D u_k. \quad (4)$$

As in the previous part, assume that the matrices  $A_k$  are always diagonal.

As in part (a) where  $\mathbf{x}_0 = 0$ , express  $y_k$  as a specific convolution-like operation of the input sequence  $\{u_\ell\}_1^L$  with a position-specific kernel  $K_k = \{K_\ell\}_1^L$ . How would you order the necessary operations to compute all the  $\{y\}$  efficiently and exploit parallelism in your hardware?


- (f) **Convolution vs. Recurrent.** Now that you have seen two ways of representing SSM (the recurrent one and the convolution one). In this question, we explore using SSM as an autoregressive generative model.  $\{x_1, \dots, x_t, \dots\}$  is a sequence of text tokens and you would like the SSM to output  $\{y_1, \dots, y_t, \dots\}$  such that  $y_{t+1} = x_t$ . **Would you use the convolution or the recurrent computation for training? What about during evaluation time where you would like to generate the text token by token?**

## 2. Attention Mechanisms for Sequence Modelling

Sequence-to-Sequence is a powerful paradigm of formulating machine learning problems. Broadly, as long as we can formulate a problem as a mapping from a sequence of inputs to a sequence of outputs, we can use sequence-to-sequence models to solve it. For example, in machine translation, we can formulate the problem as a mapping from a sequence of words in one language to a sequence of words in another language. While some RNN architectures we previously covered possess the capability to maintain a memory of the previous inputs/outputs, to compute output, the memory states need to encompass information of many previous states, which can be difficult especially when performing tasks with long-term dependencies.

To understand the limitations of vanilla RNN architectures, we consider the task of changing the case of a sentence, given a prompt token. For example, given a mixed case sequence like “ I am a student”, the model should identify this as an upper-case task based on token , and convert it to “I AM A STUDENT”. Similarly, given “ I am a student”, the lower-case task is to convert it to “i am a student”.

We can formulate this task as a character-level sequence-to-sequence problem, where the input sequence is the mixed case sentence, and the output sequence is the desired case sentence. In this exercise, we use an encoder-decoder architecture to solve the task. The encoder is a vanilla RNN that takes the input sequence as input, and outputs a sequence of hidden states. The decoder is also a vanilla RNN that takes the last hidden state from the encoder as input, and outputs the desired case sentence.



**Figure 1:** String Manipulation as a Sequence-to-Sequence Problem

- (a) What information do RNNs store?

It is important to understand how information propagates through RNNs. Particularly in the context of sequence-to-sequence models, we want to understand what information is stored in the hidden states, and what information is stored in the weights (encoder & decoder). To understand this, we consider the different components of the RNN architecture.

- **Input sequence:** The input sequence is a sequence of  $T$  tokens.


- **Encoder Weights:** The shared learnable parameters of the encoder,  $W_{\text{enc}}$
- **Bottleneck Activations:** The encoder hidden state at time  $T$ , that is passed to decoder.
- **Output sequence:** The output sequence is a sequence of  $T$  vectors (might be different length).
- **Decoder Weights:** The shared learnable parameters of the decoder,  $W_{\text{dec}}$

Consider the following questions in the context of these modules:

- i. Which of these components change during inference?
- ii. When performing gradient based updates, how are the decoder weights trained? How is gradient propagated through the encoder?
- iii. During training, what is the role of the input/output sequence?

### (b) Information Bottleneck in RNNs

Consider the architecture shown in Figure 1. This is a simple encoder-decoder architecture with a single hidden layer in the encoder and decoder. The encoder takes the input sequence as input, and outputs a sequence of hidden states. The decoder takes the last hidden state from the encoder as input, and outputs the desired case sentence. **What information needs to be stored in the hidden state to perform the upper-case/lower-case task? Are there any limitations to this architecture?**



**Figure 2:** Attention Mechanism for Sequence Modelling with RNNs.

### (c) Attention & RNNs

How does adding attention allow the model to bypass the information bottleneck? In particular, what information in the following modules would allow the model to perform the capitalization task ?


- **Encoder Weights**
- **Attention Scores**
- **Bottleneck Activations**
- **Decoder Weights**

### 3. Query-Key-Value Mechanics in Self-Attention

Self-attention is the core building block for the Transformer model, which has kickstarted the amazing progress in recent deep learning foundation models. The concept of self-attention is not restricted to Transformer exclusively though. In this question, you will be studying the detailed mechanics of the Query-Key-Value interaction in a self-attention block, in the context of RNN. Here we will be studying the case where the encoder only takes in 3 inputs, with the variables defined as:

$$h_1 = \begin{bmatrix} 5 \\ 1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 4 \\ 3 \end{bmatrix} \quad h_3 = \begin{bmatrix} -2 \\ 6 \end{bmatrix} \quad W_q = \begin{bmatrix} 1 & 0 \\ 2 & 9 \end{bmatrix} \quad W_k = \begin{bmatrix} 0 & 1 \\ 6 & 0 \end{bmatrix} \quad W_v = \begin{bmatrix} 4 & 3 \\ 9 & 1 \end{bmatrix}$$

where  $h_i$  denotes the hidden states at timestep  $i$  at some arbitrary self-attention layer. Each  $q_i, k_i, v_i$  is determined by  $q_i = W_q h_i, k_i = W_k h_i, v_i = W_v h_i$ .



**Figure 3:** Self-attention of a timestep in Encoder

- (a) **Compute**  $\hat{\alpha}_{2,1}, \hat{\alpha}_{2,2}, \hat{\alpha}_{2,3}$ .
- (b) Fig 3 shows a rough sketch of how self-attention is performed in one timestep of a Transformer block. **Write out these operations in equations, ie.  $h'_2 = \text{SelfAttention}(h_1, h_2, h_3)$ , what is `SelfAttention`?** You can define intermediate variables instead of expressing everything in one line.
- (c) For simplicity, let's use **argmax** instead of the softmax layer in the diagram. What is  $h'_2$  in this case?
- (d) In practice, it is common to have multiple self-attention operations happening in one layer. Each self-attention block is referred as a *head*, and thus the entire block is typically called *Multi-head Self-Attention (MSA)* in papers. **What's the benefit of having multiple heads?** (*Hint:* why do we want multiple kernel filters in ConvNets?)

**Contributors:**

- Naman Jain.
- Qiyang Li.
- Anant Sahai.
- Sultan Daniels.
- Gireeja Ranade.
- Kumar Krishna Agrawal.
- Kevin Li.