

1. Attention Mechanisms for Sequence Modelling

Recall the sentence capitalization problem from last week's discussion. Given an input which is a mixed case sequence like "<U> I am a student", the model should identify this as an upper-case task based on token <U>, and convert it to "I AM A STUDENT". Similarly, given "<L> I am a student", the lower-case task is to convert it to "i am a student".

This character-level task is particularly difficult for RNNs. Last week, we saw that in an encoder-decoder architecture using vanilla RNNs, all information must be encoded in the network hidden states. In particular, this results in a *information bottleneck*: the final hidden state of the encoder needs to store all the information about the input sequence, which can result in issues for long sequences.




Figure 1: Attention Mechanism for Sequence Modelling with RNNs.

Instead of storing all the information in the hidden state, we can use attention to selectively store information. **How would you modify the encoder-decoder architecture to incorporate attention, and how does this help overcome the information bottleneck?**

2. Relative Positional Embedding

Year	Model Name	Type of Positional Embedding (PE)
2018, 2019, 2020	GPT-1, 2, and 3	Learned absolute PE
2021	GOPHER	Learned relative PE
2022	PaLM	RoPE
2023, 2024	LLaMA, LLaMA-2, and LLaMA-3	RoPE
2024	Grok-1, 2	RoPE
2024, 2025	Deepseek-v2, v3	decoupled RoPE
2025	LLaMA-4	iRoPE (interleaving RoPE and NoPE)
2025	Qwen-3	RoPE
2025	Gemma 3	RoPE

Table 1: Types of positional embeddings used in large language models over the years.

- (a) **Why do we need positional encoding? Describe a situation where word order information is necessary for the task performed.**
- (b) **How does relative positional embedding work? How is it different from the absolute positional embedding?**
- (c) Consider a list of 2-dimensional vectors $\left[\left(x_1^{(1)}, x_1^{(2)} \right), \left(x_2^{(1)}, x_2^{(2)} \right), \left(x_3^{(1)}, x_3^{(2)} \right), \dots \right]$. Now for any choice of parameter ω , the corresponding RoPE encoding is

$$\text{RoPE} \left(x_t^{(1)}, x_t^{(2)}, t \right) = \begin{pmatrix} \cos t\omega & -\sin t\omega \\ \sin t\omega & \cos t\omega \end{pmatrix} \begin{pmatrix} x_t^{(1)} \\ x_t^{(2)} \end{pmatrix} = \begin{pmatrix} x_t^{(1)} \cos t\omega - x_t^{(2)} \sin t\omega \\ x_t^{(2)} \cos t\omega + x_t^{(1)} \sin t\omega \end{pmatrix}$$

Equivalently, if we write the 2-dimensional vectors as complex numbers $z_t := x_t^{(1)} + jx_t^{(2)}$, then RoPE encoding is just multiplication by a (real-valued!) parameter ω :

$$\text{RoPE}(z_t, t) = e^{jt\omega} z_t$$

The benefit of RoPE is that the dot-product between two vectors depends on their relative location only:

$$\text{RoPE}(x, m)^T \text{RoPE}(y, n) = \text{RoPE}(x, m + k)^T \text{RoPE}(y, n + k)$$

for any integer k . This is obvious if we look at RoPE in the complex number form.

Extend RoPE to work for $2n$ -dimensional vectors.

- (d) **Does adding RoPE change the number of learnable parameters in the model? If so, how many new parameters are added (compared to a similar model with no positional encoding)?**

3. NoPE (No Positional Encoding)

In models with a decoder-only architecture (such as GPT) using causal attention, it might be possible for the attention layer itself to learn to encode positional information.

Suppose we have a sequence of tokens $\{<\text{beg}>, x_2, \dots, x_n\}$ where $<\text{beg}>$ is a special token indicating the beginning of the sequence. Additionally, suppose we learn an embedding where $<\text{beg}>$ is represented by d -dimensional vector $[1, 1, e_{b,3}, \dots, e_{b,d}]^T$ and each other token x_i is represented by $[1, 0, e_{i,3}, \dots, e_{i,d}]^T$ for $i = 1, 2, \dots, n$. Here, $e_{b,j}$ and $e_{i,j}$ can be arbitrary when $j \geq 3$.

Additionally, suppose we have an attention head with query, key, value, and output matrices $W^Q \in \mathbb{R}^{d \times d}$, $W^K \in \mathbb{R}^{d \times d}$, $W^V \in \mathbb{R}^{d \times d}$ with the following learned values:

$$W^K = \begin{bmatrix} 1 & 1 & \cdots & 1 \\ 0 & 0 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 0 \end{bmatrix}, \quad W^V = \begin{bmatrix} 0 & 0 & \cdots & 0 \\ 1 & 0 & \cdots & 0 \\ 0 & 0 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 0 \end{bmatrix}$$

and W^Q is arbitrary.

Using the convention that weight matrices are multiplied on the right of the token embeddings, **show that the output of the attention head for token x_t encodes the position t , and explain how this can pass positional information to subsequent layers in the model.**

4. Transformer Architecture




Figure 2: LLaMA architecture diagram. Many modern transformers use rotary positional embeddings (RoPE) and SwiGLU activation function (source: <https://github.com/hkproj/pytorch-llama-notes/>).

- (a) **What is the advantage of multi-headed attention?** Give some examples of linguistic structures that can be found using multi-headed attention (e.g. subject-verb agreement, translation between languages).

- (b) Let's say we're using argmax attention, which uses argmax rather than softmax, like we saw previously.
What is the size of the receptive field of a node at level n ...
If we have only a single head?
If we have two heads?
If we have h heads?
- (c) For input sequences of length M and output sequences of length N , **what are the complexities of** (1) **Non-Causal Self-Attention¹** (on the input sequence), (2) **Non-Causal Self-Attention** (on the input sequence), but with multi-query attention, (3) **Cross Attention** (from the output sequence to the input sequence), (4) **Causal Self-Attention** (on the output sequence). Let d be the hidden dimension of the network (both the encoder and the decoder part), and h be the number of heads.

Recall: Normal multi-headed attention is:

$$\text{MultiHeadedAttention}(Q, K, V) = \text{Concat}_{i \in [\# \text{ heads}]} \left(\text{Attention} \left(XW_i^Q, XW_i^K, XW_i^V \right) \right) W^O$$

For this question, we assume W_i^Q, W_i^K, W_i^V are of shape $d \times d$ and X is shape $M \times d$ where each row of X is a token embedding. We also assume that we use the naive matrix multiplication, that is, if A, B are of shape $m \times n, n \times k$, then AB takes mnk multiply-accumulate operations.

With multi-query attention, there is just one W^K, W^V , thus:

$$\text{MultiQueryAttention}(Q, K, V) = \text{Concat}_{i \in [\# \text{ heads}]} \left(\text{Attention} \left(XW_i^Q, XW^K, XW^V \right) \right) W^O$$

- (d) **How many operations does the multi-query attention save compared to multi-head attention?**
What other some other advantages?

Contributors:

- Kumar Krishna Agrawal.
- Kevin Li.
- Anant Sahai.
- Olivia Watkins.
- Jerome Quenum.
- Saagar Sanghavi.
- Yuxi Liu.
- Qiyang Li.
- CS 182/282A Staff from previous semesters.
- Lance Mathias.

¹This is the “normal” self-attention from lecture.